



Introduction

Week1_발표자: 구미진

목차


01 Intro

02 History of Computer Vision

03 Various tasks of Computer Vision



Intro



01 Intro



Lecture 1:
Introduction

모두 재생

Stanford University CS231n, Spring 2017

동영상 16개 • 조회수 894,744회 • 최종 업데이트:
2017. 8. 11.

≡+ ↺ ↻ ...

CS231n: Convolutional Neural Networks for Visual Recognition

Spring 2017

<http://cs231n.stanford.edu/>

CS231n Home


Course Notes

Coursework Schedule


Office Hours

Lecture Videos


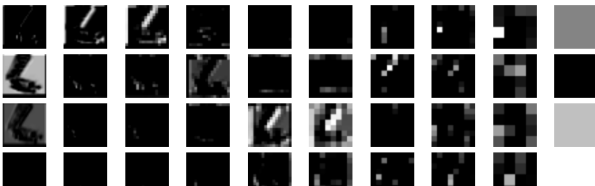
Piazza




CS231n: Convolutional Neural Networks for Visual Recognition




Stanford - Spring 2021





CS231n: Convolutional Neural Networks for Visual Recognition



Schedule and Syllabus

Unless otherwise specified the course lectures and meeting times are Tuesday and Thursday 12pm to 1:20pm in the NVIDIA Auditorium in the Huang Engineering Center. ([map](#))

This is the syllabus for the **Spring 2017** iteration of the course. The syllabus for the [Winter 2016](#) and [Winter 2015](#) iterations of this course are still available.

Event Type	Date	Description	Course Materials
Lecture 1	Tuesday April 4	Course Introduction Computer vision overview Historical context Course logistics	[slides] [video]
Lecture 2	Thursday April 6	Image Classification The data-driven approach K-nearest neighbor Linear classification I	[slides] [video] [python/numpy tutorial] [image classification notes] [linear classification notes]

CS231n Convolutional Neural Networks for Visual Recognition

Course Website

These notes accompany the Stanford CS class [CS231n: Convolutional Neural Networks for Visual Recognition](#). For questions/concerns/bug reports, please submit a pull request directly to our [git repo](#).

Spring 2021 Assignments

Assignment #1: Image Classification, kNN, SVM, Softmax, Fully Connected Neural Network

Assignment #2: Fully Connected and Convolutional Nets, Batch Normalization, Dropout, Frameworks

Assignment #3: Image Captioning with RNNs and Transformers, Network Visualization, Generative Adversarial Networks, Self-Supervised Contrastive Learning

Module 0: Preparation

Software Setup

Python / Numpy Tutorial (with Jupyter and Colab)

Google Cloud Tutorial

Module 1: Neural Networks

[Image Classification: Data-driven Approach, k-Nearest Neighbor, train/val/test splits, L1/L2 distances, hyperparameter search, cross-validation](#)

Linear classification: Support Vector Machine, Softmax
[parameteric approach, bias trick, hinge loss, cross-entropy loss, L2 regularization, web demo](#)

Optimization: Stochastic Gradient Descent
[optimization landscapes, local search, learning rate, analytic/numerical gradient](#)

Cs231n은 컴퓨터 비전에 관한 수업으로,
Neural network(신경망)과 CNN(Deep Learning)과 관련된 부분을 중점적으로 공부합니다.

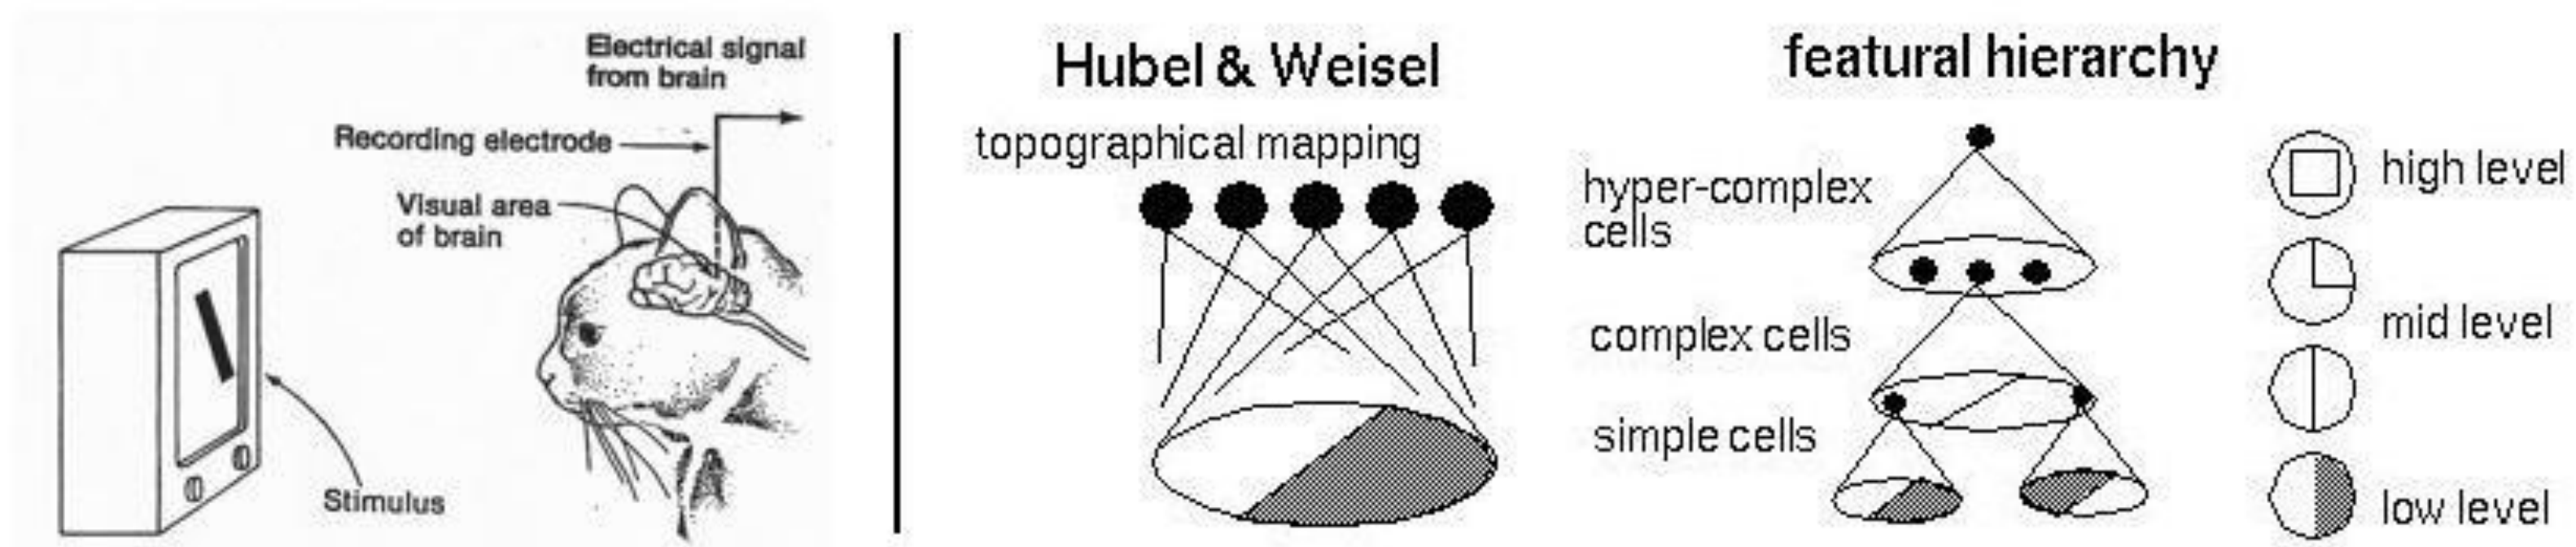
History of Computer Vision



02 동물의 시각 처리 메커니즘

Hubel & Wiesel, 1959

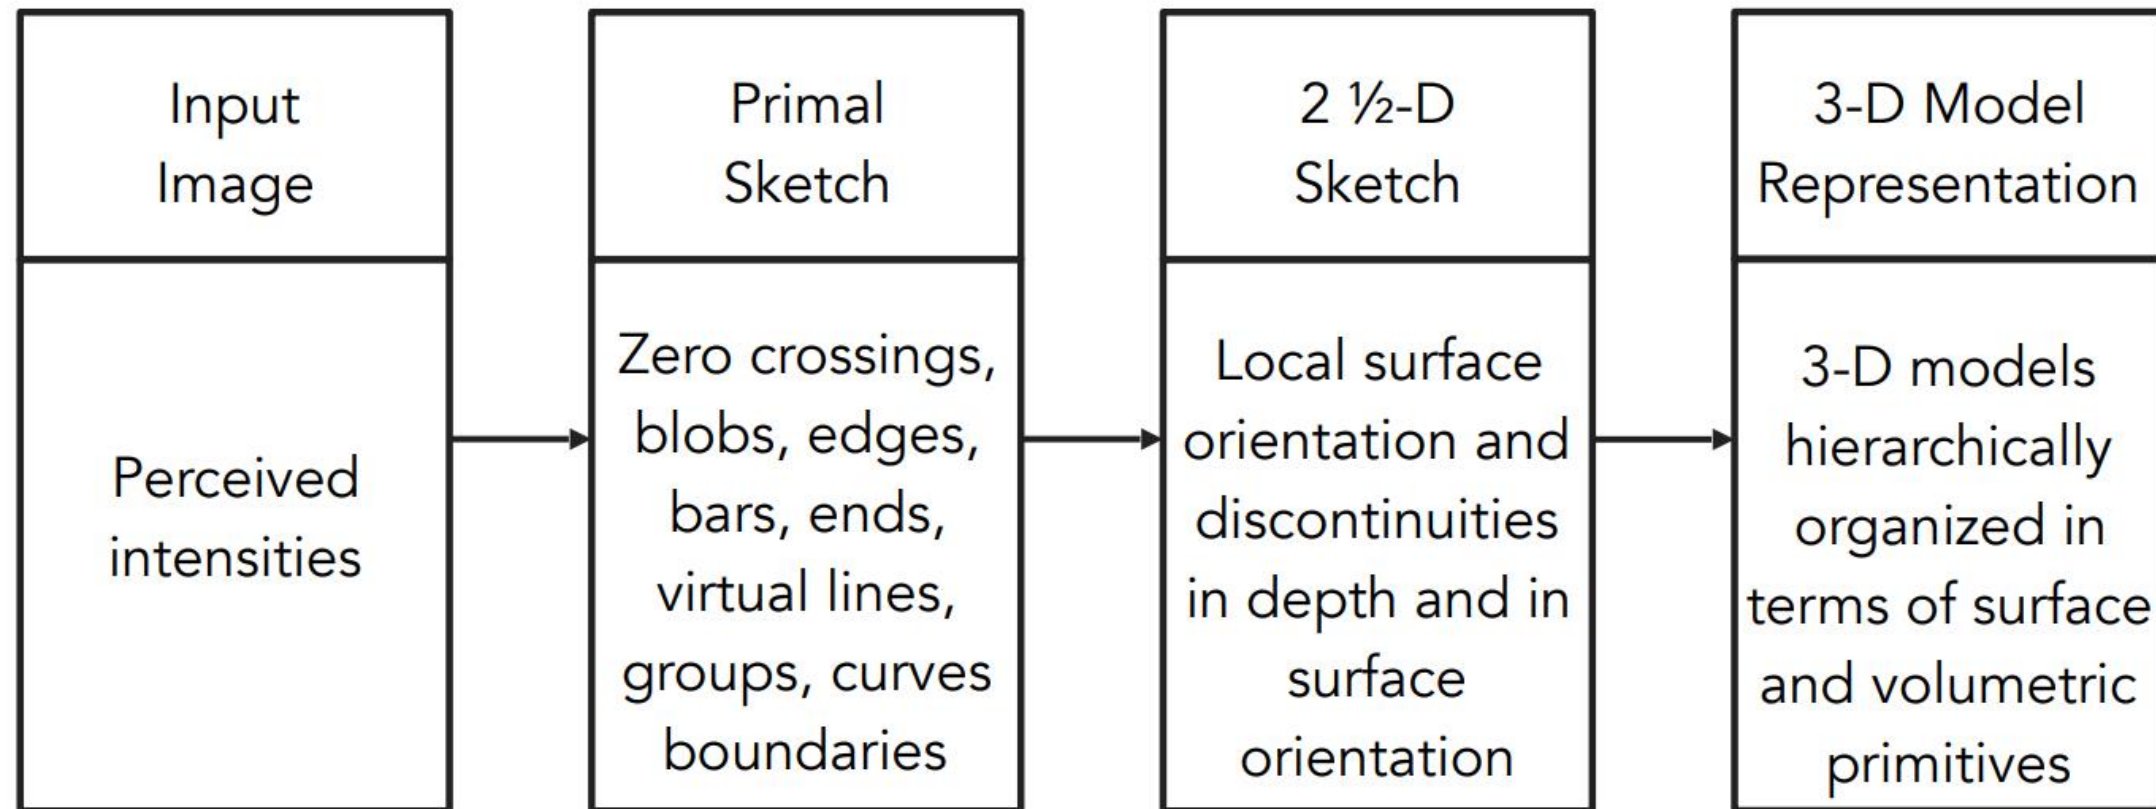
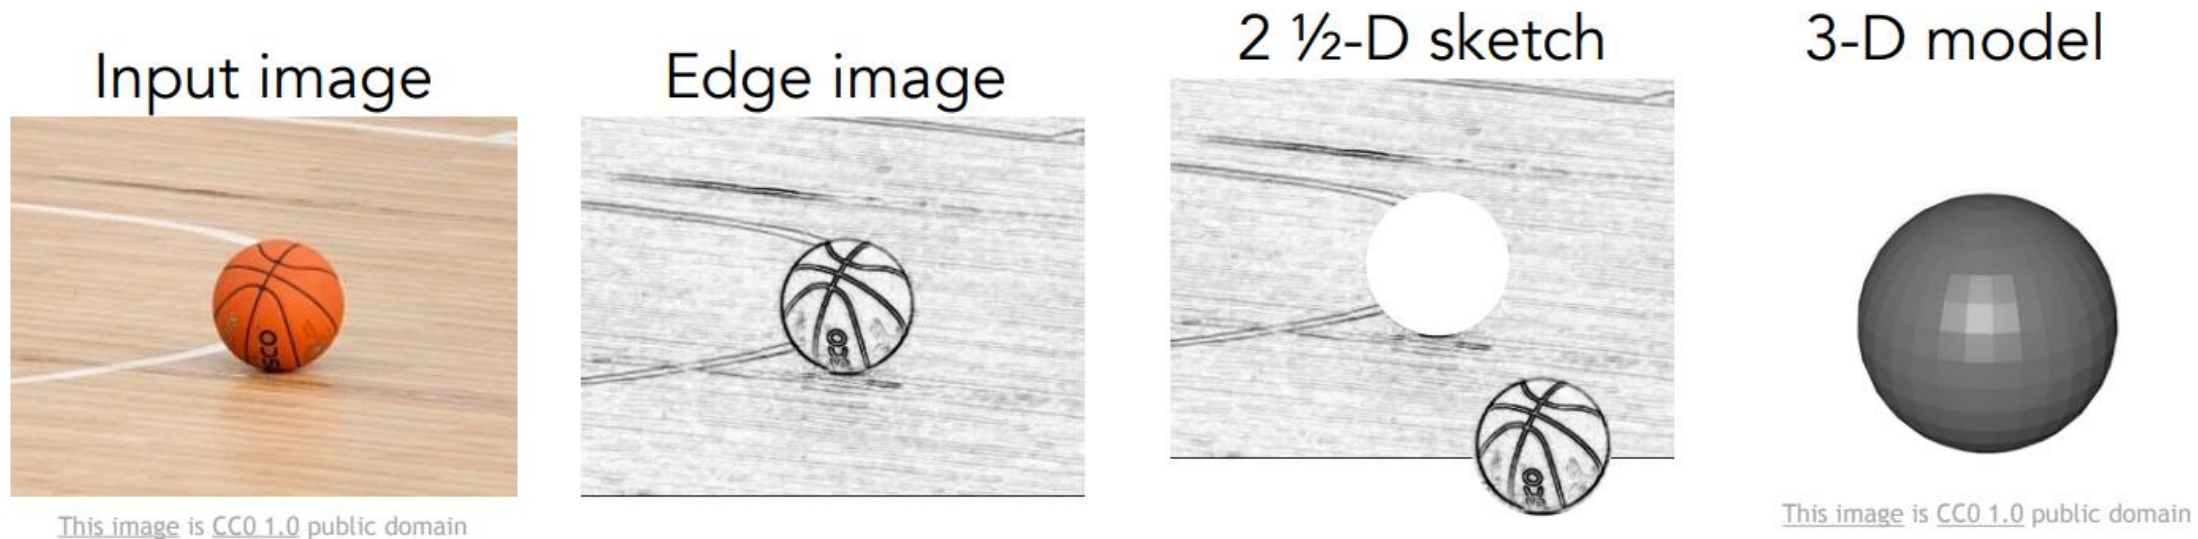
동물의 시각적 처리 메커니즘은 무엇일까?



“시각 처리는 edges와 같은 단순한 구조로 시작하고, 그 정보가 layer를 거치면서 점점 복잡해져 실제 이미지를 시각적으로 인식할 수 있게 된다.”

02 시각적 표현의 단계

눈으로 Input 이미지를 보고 3D 표현으로 재구성하려면 몇 단계 과정을 거쳐야 함



Stages of Visual Representation, David Marr, 1970s

02 Object Segmentation

Normalized cut (Shi & Malik, 1997)


객체 분할(object segmentation) : 이미지의 각 픽셀을 의미 있는 영역들로 grouping하는 방법



02 Object recognition

SIFT(Scale-Invariant Feature Transform)

- 이미지의 크기와 회전에 불변하는 특징을 추출하는 특징 기반 객체 인식 알고리즘
- 서로 다른 두 이미지에서 feature를 각각 추출한 다음, 다른 이미지에 추출한 특징들을 매칭시키는 방법



02 Image classification - ILSVRC


IMAGENET Large Scale Visual Recognition Challenge

The Image Classification Challenge:

1,000 object classes

1,431,167 images

Overfitting을 극복하면서,
세상의 모든 객체를 인식 해보자!



02 Image classification - ILSVRC


IMAGENET Large Scale Visual Recognition Challenge

The Image Classification Challenge:

1,000 object classes

1,431,167 images


Overfitting을 극복하면서,
세상의 모든 객체를 인식 해보자!




CNN 기반의 AlexNet

02 CNN

LeNet (1998)




AlexNet (2012)




02 CNN

LeNet (1998)




AlexNet (2012)



1. 연산량 증가 (컴퓨터의 계산속도, GPU)
2. 데이터의 증가

Various tasks of Computer Vision



03 Various tasks of Computer Vision

[Browse SoTA](#) > Computer Vision

Computer Vision

2667 benchmarks • 949 tasks • 1755 datasets • 23343 papers with code

Image Classification





Image Classification

📊 299 benchmarks


2167 papers with code



Knowledge Distillation

📊 3 benchmarks


456 papers with code



Few-Shot Image Classification

📊 80 benchmarks


121 papers with code



Fine-Grained Image Classification

📊 32 benchmarks

94 papers with code




Semi-Supervised Image Classification

📊 34 benchmarks

81 papers with code

▶ [See all 22 tasks](#)


Semantic Segmentation



Semantic Segmentation

📊 128 benchmarks


2499 papers with code



Tumor Segmentation

📊 1 benchmark


115 papers with code



3D Semantic Segmentation

📊 8 benchmarks


78 papers with code



Panoptic Segmentation

📊 10 benchmarks

77 papers with code



Scene Segmentation

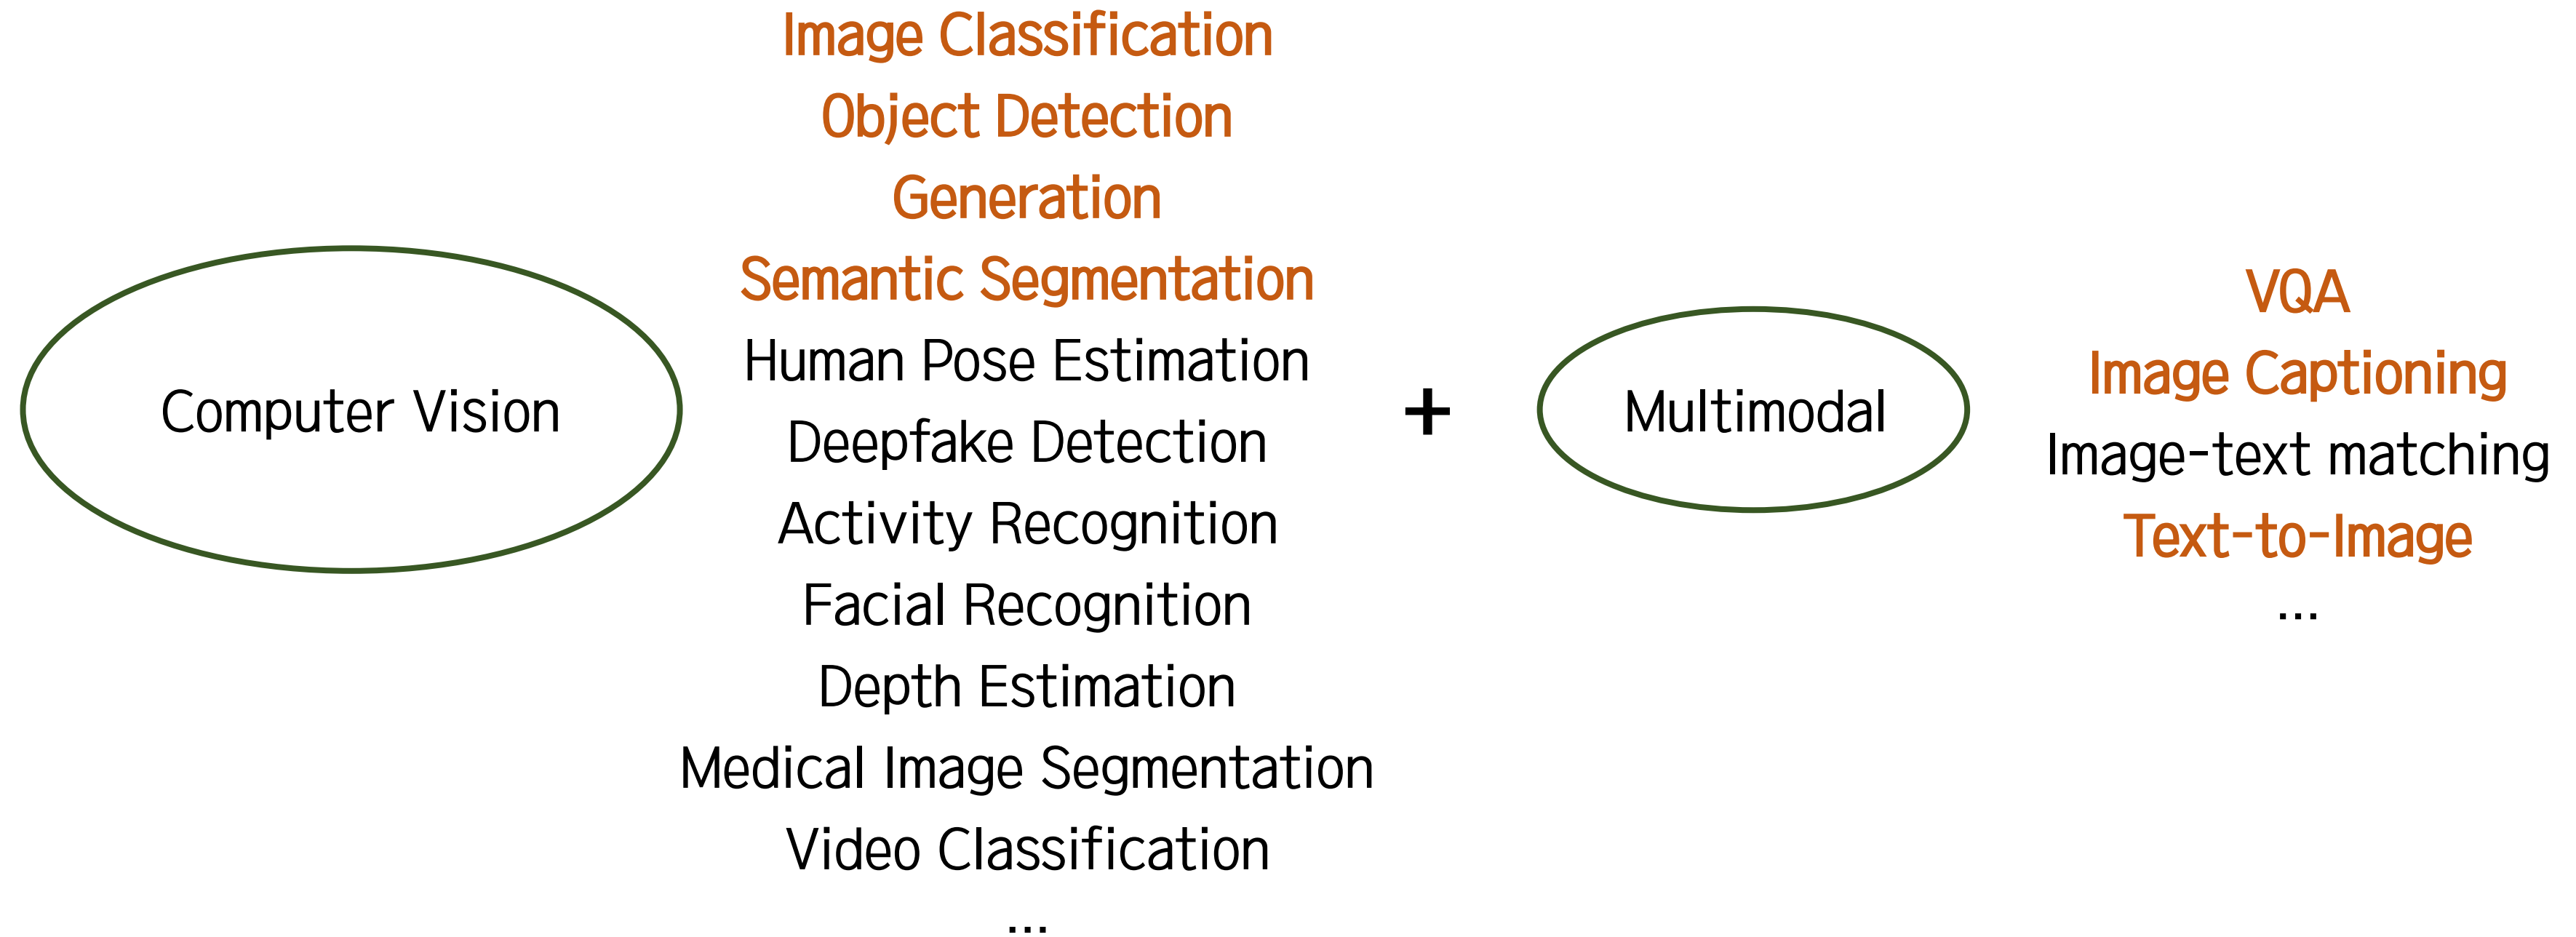
📊 6 benchmarks

66 papers with code

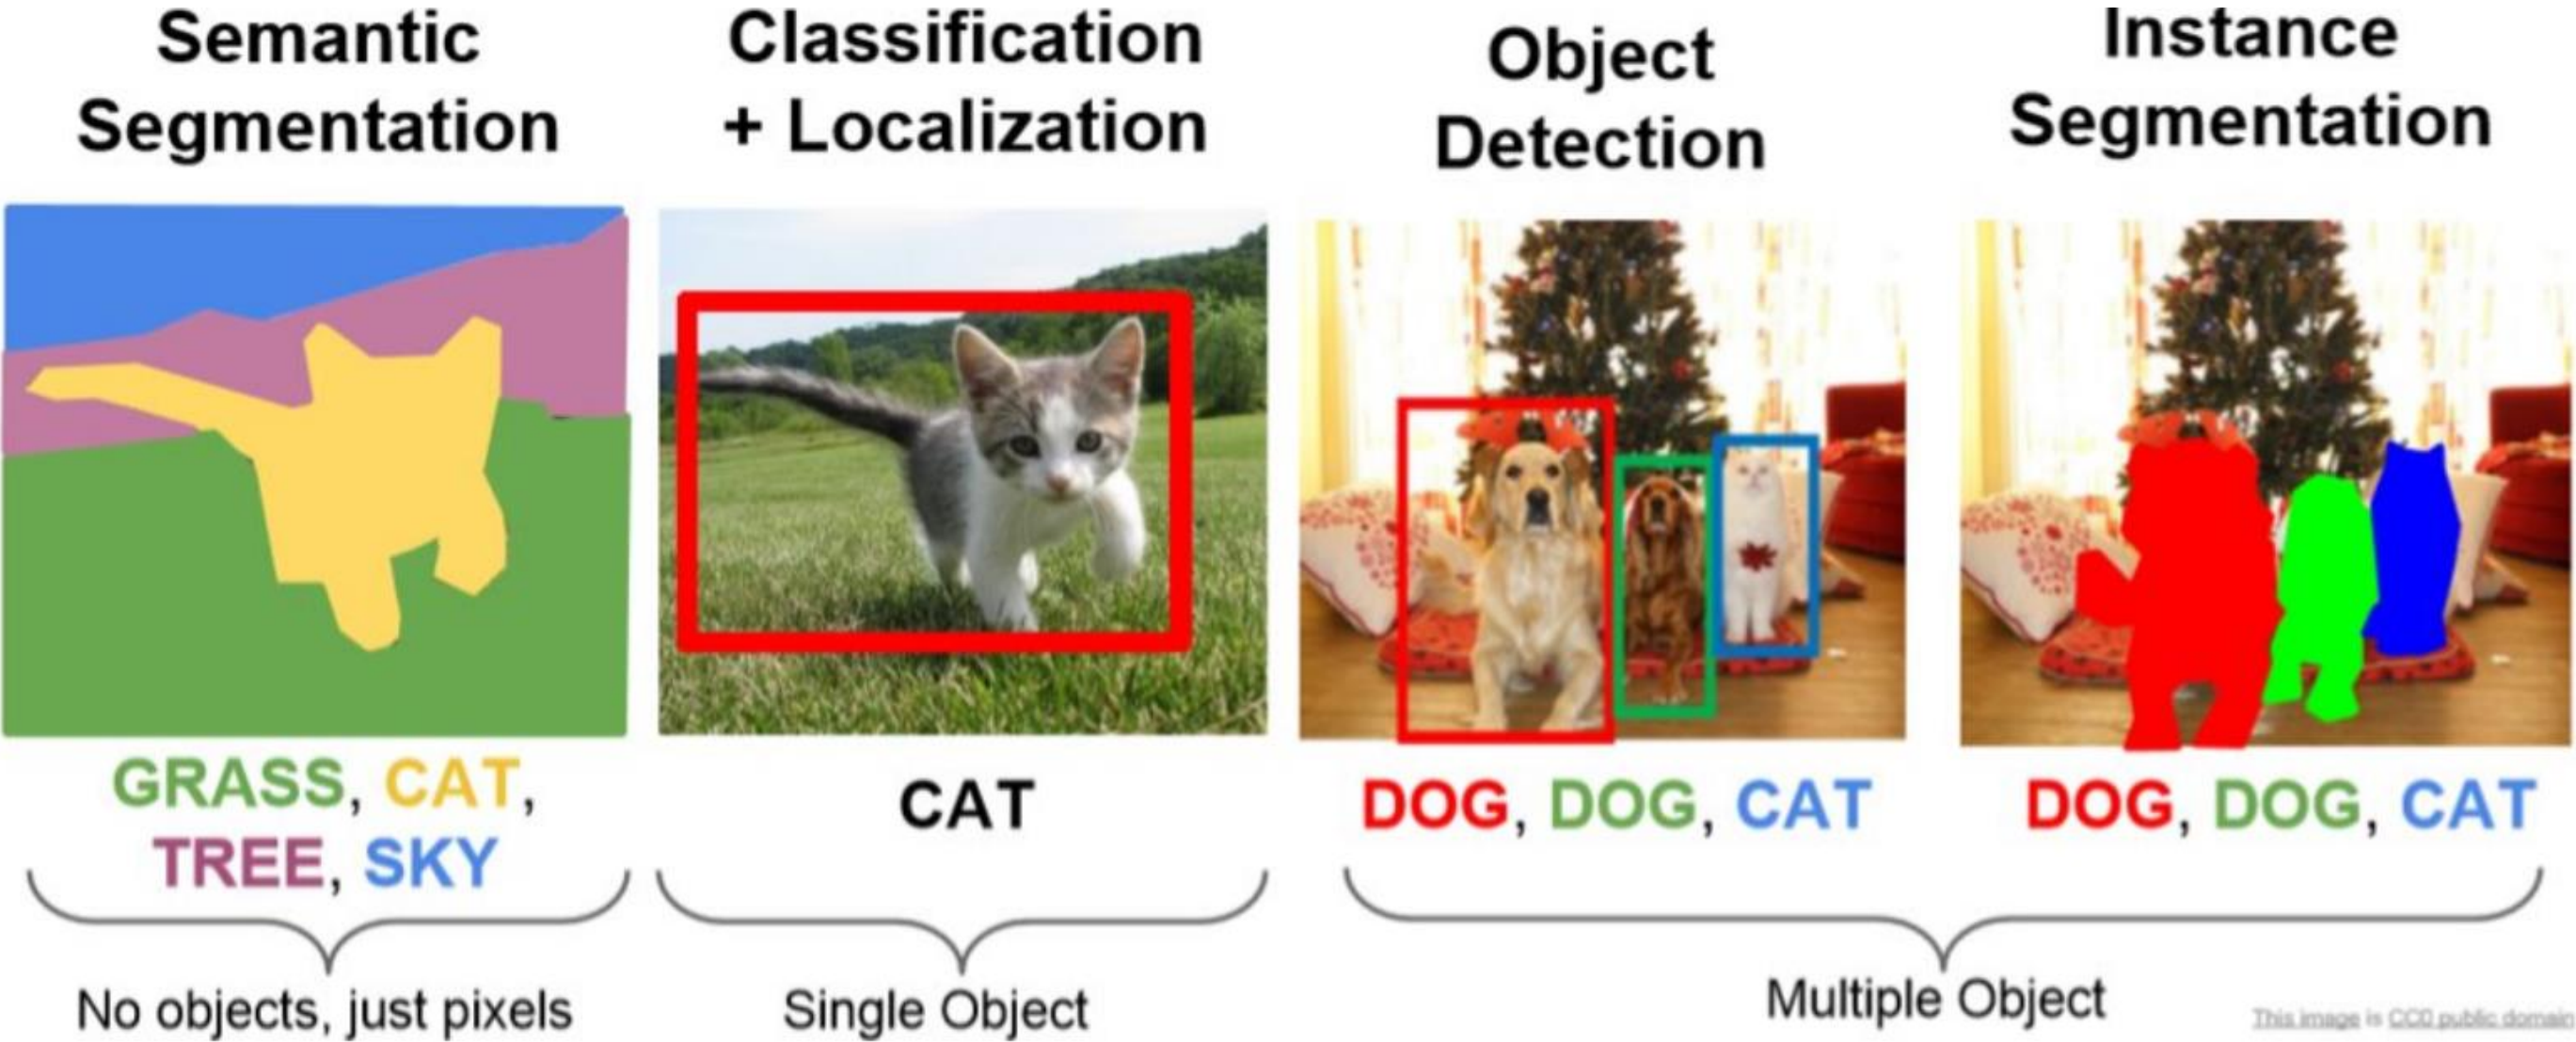
▶ [See all 19 tasks](#)

Object Detection

03 Various tasks of Computer Vision



03 Various tasks of Computer Vision

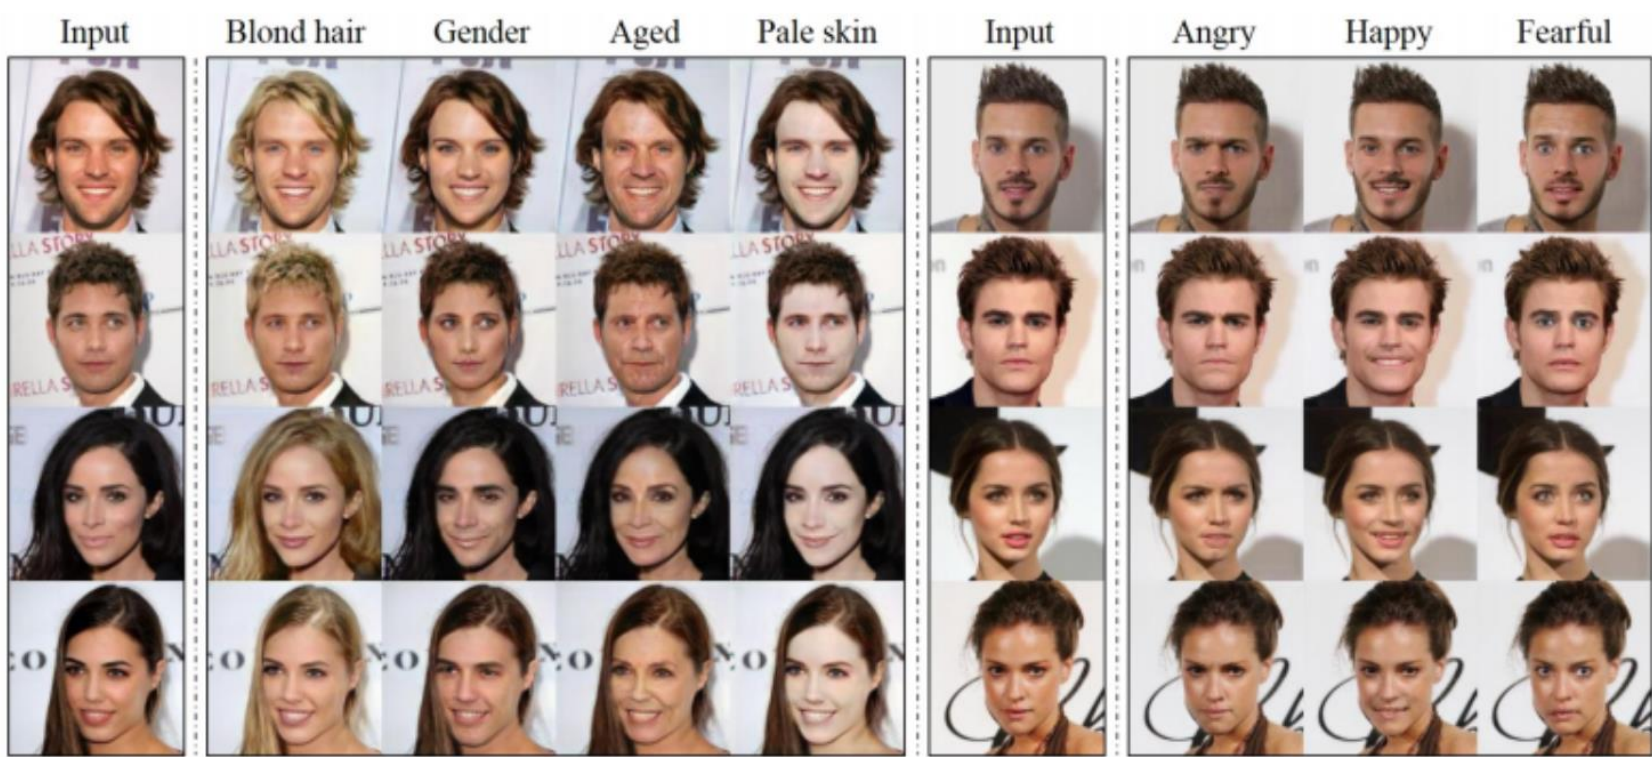


This image is CC0 public domain

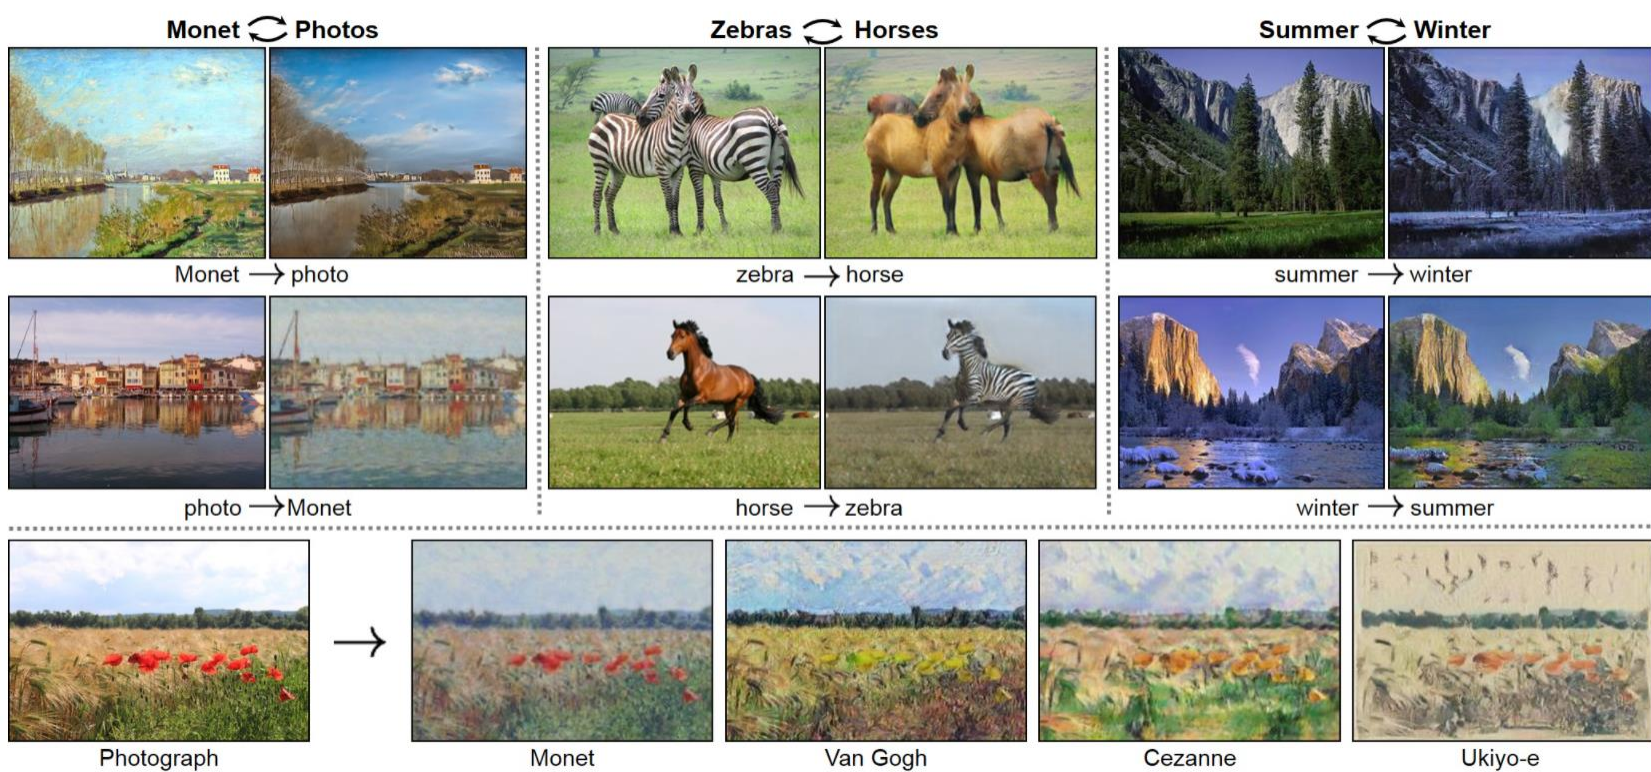
03 Various tasks of Computer Vision

GAN(Generative Adversarial Network)

StarGan



CycleGan



03 Various tasks of Computer Vision

Image Captioning

<p>A young boy is playing basketball.</p> 	<p>Two dogs play in the grass.</p> 	<p>A dog swims in the water.</p> 	<p>A little girl in a pink shirt is swinging.</p> 
<p>A group of people walking down a street.</p> 	<p>A group of women dressed in formal attire.</p> 	<p>Two children play in the water.</p> 	<p>A dog jumps over a hurdle.</p> 

VQA

<p>Who is wearing glasses?</p> <p>man</p> 	<p>woman</p> 	<p>Where is the child sitting?</p> <p>fridge</p> 	<p>arms</p> 
<p>Is the umbrella upside down?</p> <p>yes</p> 	<p>no</p> 	<p>How many children are in the bed?</p> <p>2</p> 	<p>1</p> 

THANK YOU

