



Deep Reinforcement Learning

Week16 구미진, 민소연

Index

01 Reinforcement Learning

02 Markov Decision Processes

03 Q-Learning

04 Policy Gradients



Reinforcement Learning



01 Supervised vs Unsupervised Learning

Data: (x, y)

x is data, y is label

Goal: Learn a *function* to map $x \rightarrow y$

Examples: Classification, regression, object detection, semantic segmentation, image captioning, etc.



→ Cat

Classification

Data: x

Just data, no labels!



Goal: Learn some underlying hidden *structure* of the data

Examples: Clustering, dimensionality reduction, feature learning, density estimation, etc.



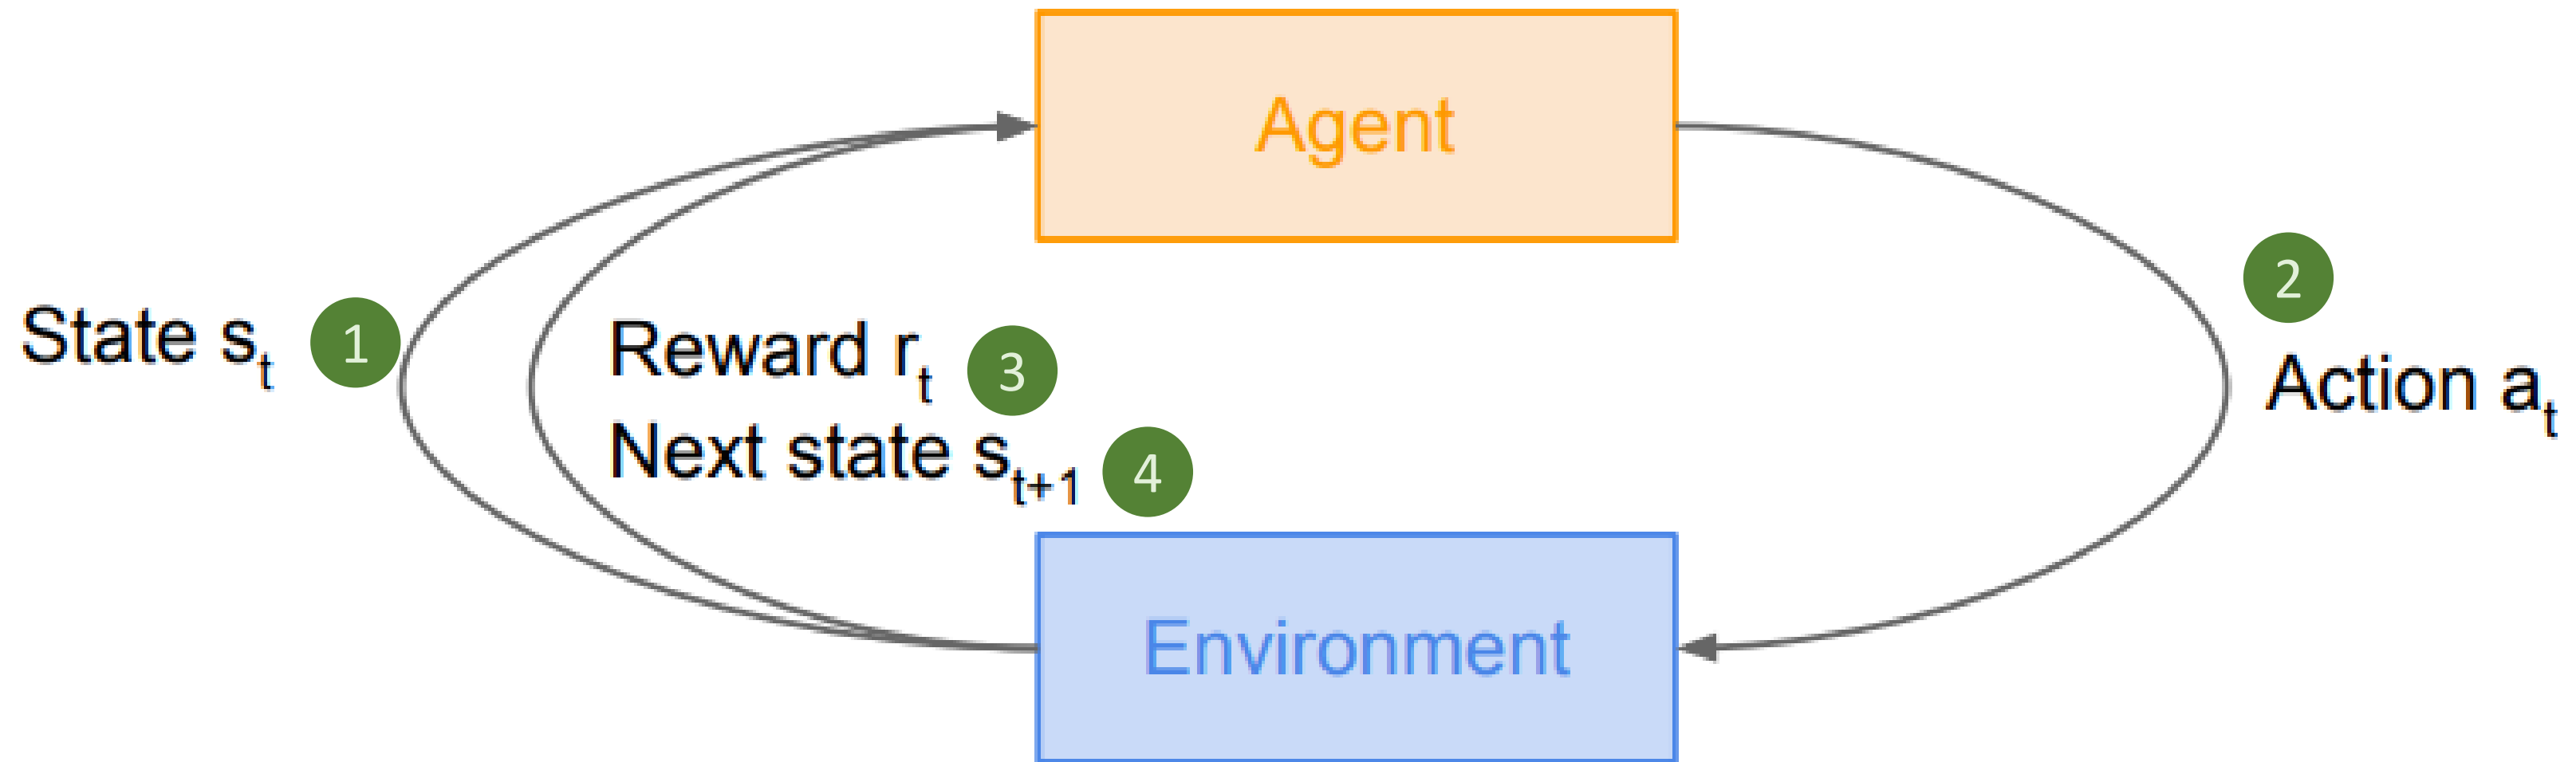
Figure copyright Ian Goodfellow, 2016. Reproduced with permission.

1-d density estimation




2-d density estimation

02 Reinforcement Learning



03 Examples




Objective: Balance a pole on top of a movable cart

State: angle, angular speed, position, horizontal velocity

Action: horizontal force applied on the cart

Reward: 1 at each time step if the pole is upright



Objective: Make the robot move forward

State: Angle and position of the joints

Action: Torques applied on joints

Reward: 1 at each time step upright + forward movement

Atari Games




Objective: Complete the game with the highest score

State: Raw pixel inputs of the game state

Action: Game controls e.g. Left, Right, Up, Down

Reward: Score increase/decrease at each time step

Go



Objective: Win the game!

State: Position of all pieces

Action: Where to put the next piece down

Reward: 1 if win at the end of the game, 0 otherwise

Markov Decision Process



01 Markov Decision Process

Markov Decision Process

- Mathematical formulation of the RL problem
- **Markov property**: Current state completely characterises the state of the world

Defined by: $(\mathcal{S}, \mathcal{A}, \mathcal{R}, \mathbb{P}, \gamma)$

\mathcal{S} : set of possible states

\mathcal{A} : set of possible actions

\mathcal{R} : distribution of reward given (state, action) pair

\mathbb{P} : transition probability i.e. distribution over next state given (state, action) pair

γ : discount factor

- A policy π is a function from \mathcal{S} to \mathcal{A} that specifies what action to take in each state
- **Objective**: find policy π^* that maximizes cumulative discounted reward: $\sum_{t \geq 0} \gamma^t r_t$

02 A simple MDP

actions = {

1. right →


2. left ←

3. up ↑

4. down ↓


}

states




Set a negative “reward”
for each transition
(e.g. $r = -1$)

Objective: reach one of terminal states (greyed out) in
least number of actions



Random Policy



Optimal Policy

03 How to Find Optimal Policy

The optimal policy π^*

We want to find optimal policy π^* that maximizes the sum of rewards.

How do we handle the randomness (initial state, transition probability...)?

Maximize the **expected sum of rewards!**

Formally: $\pi^* = \arg \max_{\pi} \mathbb{E} \left[\sum_{t \geq 0} \gamma^t r_t | \pi \right]$ with $s_0 \sim p(s_0)$, $a_t \sim \pi(\cdot | s_t)$, $s_{t+1} \sim p(\cdot | s_t, a_t)$

Q-Learning



01 Value function and Q-value function

Definitions: Value function and Q-value function

Following a policy produces sample trajectories (or paths) $s_0, a_0, r_0, s_1, a_1, r_1, \dots$

How good is a state?

The **value function** at state s , is the expected cumulative reward from following the policy from state s :

$$V^\pi(s) = \mathbb{E} \left[\sum_{t \geq 0} \gamma^t r_t \mid s_0 = s, \pi \right]$$

How good is a state-action pair?

The **Q-value function** at state s and action a , is the expected cumulative reward from taking action a in state s and then following the policy:

$$Q^\pi(s, a) = \mathbb{E} \left[\sum_{t \geq 0} \gamma^t r_t \mid s_0 = s, a_0 = a, \pi \right]$$

Bellman equation

The optimal Q-value function Q^* is the maximum expected cumulative reward achievable from a given (state, action) pair:

$$Q^*(s, a) = \max_{\pi} \mathbb{E} \left[\sum_{t \geq 0} \gamma^t r_t \mid s_0 = s, a_0 = a, \pi \right]$$

Q^* satisfies the following **Bellman equation**:

$$Q^*(s, a) = \mathbb{E}_{s' \sim \mathcal{E}} \left[r + \gamma \max_{a'} Q^*(s', a') \mid s, a \right]$$

Intuition: if the optimal state-action values for the next time-step $Q^*(s', a')$ are known, then the optimal strategy is to take the action that maximizes the expected value of $r + \gamma Q^*(s', a')$

The optimal policy π^* corresponds to taking the best action in any state as specified by Q^*

Solving for the optimal policy

Value iteration algorithm: Use Bellman equation as an iterative update

$$Q_{i+1}(s, a) = \mathbb{E} \left[r + \gamma \max_{a'} Q_i(s', a') | s, a \right]$$

Q_i will converge to Q^* as $i \rightarrow \infty$

What's the problem with this?

Not scalable. Must compute $Q(s, a)$ for every state-action pair. If state is e.g. current game state pixels, computationally infeasible to compute for entire state space!

Solution: use a function approximator to estimate $Q(s, a)$. E.g. a neural network!

Solving for the optimal policy: Q-learning

Q-learning: Use a function approximator to estimate the action-value function

$$Q(s, a; \theta) \approx Q^*(s, a)$$

function parameters (weights)

If the function approximator is a deep neural network => **deep q-learning!**

04 Q-Learning

Remember: want to find a Q-function that satisfies the Bellman Equation:

$$Q^*(s, a) = \mathbb{E}_{s' \sim \mathcal{E}} \left[r + \gamma \max_{a'} Q^*(s', a') | s, a \right]$$

Forward Pass

Loss function: $L_i(\theta_i) = \mathbb{E}_{s, a \sim \rho(\cdot)} \left[(y_i - Q(s, a; \theta_i))^2 \right]$

where $y_i = \mathbb{E}_{s' \sim \mathcal{E}} \left[r + \gamma \max_{a'} Q(s', a'; \theta_{i-1}) | s, a \right]$

Iteratively try to make the Q-value close to the target value (y_i) it should have, if Q-function corresponds to optimal Q^* (and optimal policy π^*)


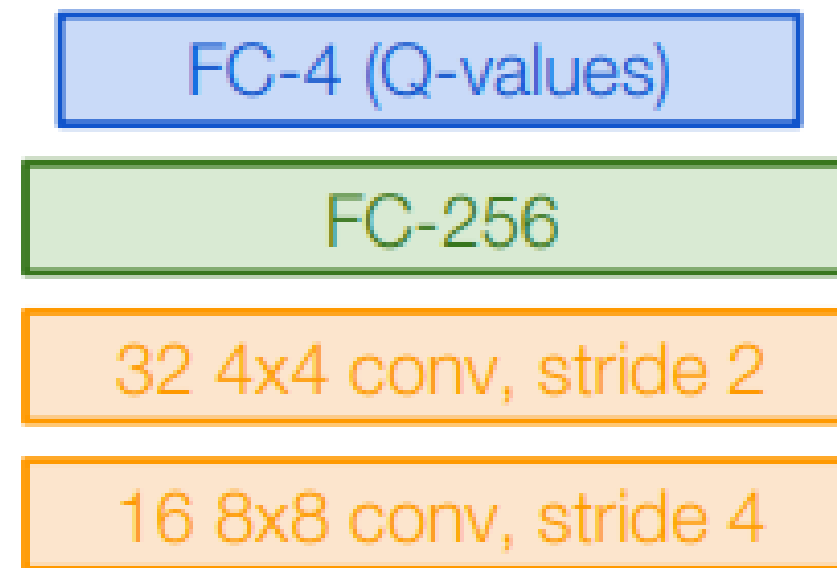
Backward Pass

Gradient update (with respect to Q-function parameters θ):

$$\nabla_{\theta_i} L_i(\theta_i) = \mathbb{E}_{s, a \sim \rho(\cdot); s' \sim \mathcal{E}} \left[r + \gamma \max_{a'} Q(s', a'; \theta_{i-1}) - Q(s, a; \theta_i) \right] \nabla_{\theta_i} Q(s, a; \theta_i)$$

Q-network Architecture

$Q(s, a; \theta)$:
neural network
with weights θ



Last FC Layer has 4-d output
(if 4 actions), corresponding to
 $Q(s_t, a_1)$, $Q(s_t, a_2)$, $Q(s_t, a_3)$, $Q(s_t, a_4)$

Familiar conv layers, FC layer

Input: State s_t

Current state s_t : 84x84x4 stack of last 4 frames
(after RGB->grayscale conversion, downsampling, and cropping)

Training the Q-network: Experience Replay

Learning from batches of consecutive samples is problematic:

- Samples are correlated => inefficient learning
- Current Q-network parameters determines next training samples (e.g. if maximizing action is to move left, training samples will be dominated by samples from left-hand size) => can lead to bad feedback loops

Address these problems using **experience replay**

- Continually update a **replay memory** table of transitions (s_t, a_t, r_t, s_{t+1}) as game (experience) episodes are played
- Train Q-network on random minibatches of transitions from the replay memory, instead of consecutive samples

Each transition can also contribute
to multiple weight updates
=> greater data efficiency

04 Q-Learning

Putting it together: Deep Q-Learning with Experience Replay

Algorithm 1 Deep Q-learning with Experience Replay

Initialize replay memory \mathcal{D} to capacity N

Initialize action-value function Q with random weights

} Initialize replay memory, Q-network

for episode = 1, M **do** ← Play M episodes (full games)

 Initialize sequence $s_1 = \{x_1\}$ and preprocessed sequenced $\phi_1 = \phi(s_1)$

← Initialize state (starting game screen pixels) at the beginning of each episode

for $t = 1, T$ **do**

 With probability ϵ select a random action a_t

 otherwise select $a_t = \max_a Q^*(\phi(s_t), a; \theta)$

 Execute action a_t in emulator and observe reward r_t and image x_{t+1}

 Set $s_{t+1} = s_t, a_t, x_{t+1}$ and preprocess $\phi_{t+1} = \phi(s_{t+1})$

 Store transition $(\phi_t, a_t, r_t, \phi_{t+1})$ in \mathcal{D}

 Sample random minibatch of transitions $(\phi_j, a_j, r_j, \phi_{j+1})$ from \mathcal{D}

 Set $y_j = \begin{cases} r_j & \text{for terminal } \phi_{j+1} \\ r_j + \gamma \max_{a'} Q(\phi_{j+1}, a'; \theta) & \text{for non-terminal } \phi_{j+1} \end{cases}$

 Perform a gradient descent step on $(y_j - Q(\phi_j, a_j; \theta))^2$ according to equation 3

end for

end for

04 Q-Learning

Putting it together: Deep Q-Learning with Experience Replay

Algorithm 1 Deep Q-learning with Experience Replay

Initialize replay memory \mathcal{D} to capacity N

Initialize action-value function Q with random weights

for episode = 1, M **do**

 Initialize sequence $s_1 = \{x_1\}$ and preprocessed sequenced $\phi_1 = \phi(s_1)$

for $t = 1, T$ **do**

 With probability ϵ select a random action a_t
 otherwise select $a_t = \max_a Q^*(\phi(s_t), a; \theta)$

 Execute action a_t in emulator and observe reward r_t and image x_{t+1}

 Set $s_{t+1} = s_t, a_t, x_{t+1}$ and preprocess $\phi_{t+1} = \phi(s_{t+1})$

 Store transition $(\phi_t, a_t, r_t, \phi_{t+1})$ in \mathcal{D}

 Sample random minibatch of transitions $(\phi_j, a_j, r_j, \phi_{j+1})$ from \mathcal{D}

 Set $y_j = \begin{cases} r_j & \text{for terminal } \phi_{j+1} \\ r_j + \gamma \max_{a'} Q(\phi_{j+1}, a'; \theta) & \text{for non-terminal } \phi_{j+1} \end{cases}$

 Perform a gradient descent step on $(y_j - Q(\phi_j, a_j; \theta))^2$ according to equation 3

end for

end for

For each timestep t of the game

With small probability, select a random action(explore), otherwise select greedy action from current policy

Take the action (a_t), and observe the reward r_t and next state s_{t+1}

04 Q-Learning

Putting it together: Deep Q-Learning with Experience Replay

Algorithm 1 Deep Q-learning with Experience Replay

Initialize replay memory \mathcal{D} to capacity N

Initialize action-value function Q with random weights

for episode = 1, M **do**

 Initialize sequence $s_1 = \{x_1\}$ and preprocessed sequenced $\phi_1 = \phi(s_1)$

for $t = 1, T$ **do**

 With probability ϵ select a random action a_t

 otherwise select $a_t = \max_a Q^*(\phi(s_t), a; \theta)$

 Execute action a_t in emulator and observe reward r_t and image x_{t+1}

 Set $s_{t+1} = s_t, a_t, x_{t+1}$ and preprocess $\phi_{t+1} = \phi(s_{t+1})$

 Store transition $(\phi_t, a_t, r_t, \phi_{t+1})$ in \mathcal{D}

 Sample random minibatch of transitions $(\phi_j, a_j, r_j, \phi_{j+1})$ from \mathcal{D}

 Set $y_j = \begin{cases} r_j & \text{for terminal } \phi_{j+1} \\ r_j + \gamma \max_{a'} Q(\phi_{j+1}, a'; \theta) & \text{for non-terminal } \phi_{j+1} \end{cases}$

 Perform a gradient descent step on $(y_j - Q(\phi_j, a_j; \theta))^2$ according to equation 3

end for

end for

Store transition in replay memory

Experience Replay:
Sample a random minibatch of
transitions from replay memory
and perform a gradient
descent step

Policy Gradients



01 Policy Gradients

Q-learning의 문제점

- Q-function이 아주 복잡함
 - 모든 (state, action) 쌍들을 학습해야 함
 - 로봇이 어떤 물체를 손으로 잡는 문제를 해결 한다고 했을 때는?
- > 로봇의 모든 관절의 위치와 각도가 이룰 수 있는 모든 경우의 수에 대해 모든 (state, action)을 학습시켜야 함

01 Policy Gradients

Q-learning의 문제점

- Q-function이 아주 복잡함
 - 모든 (state, action) 쌍들을 학습해야 함
 - 로봇이 어떤 물체를 손으로 잡는 문제를 해결 한다고 했을 때는?
- > 로봇의 모든 관절의 위치와 각도가 이룰 수 있는 모든 경우의 수에 대해 모든 (state, action)을 학습시켜야 함
- > 모든 (state, action)을 학습시키는 대신, **정책** 자체를 학습시키는 방법!

“Policy Gradients”

01 Policy Gradients

매개변수화된 정책들의 집합
(가중치 θ)

$$\Pi = \{\pi_{\theta}, \theta \in \mathbb{R}^m\}$$

미래에 받을 보상들의 누적 합의 기댓값

$$J(\theta) = \mathbb{E} \left[\sum_{t \geq 0} \gamma^t r_t | \pi_{\theta} \right]$$

$J(\theta)$ 를 최대로 만드는 최적의 정책 θ^*

$$\theta^* = \arg \max_{\theta} J(\theta)$$

02 REINFORCE algorithm

경로에 대한 미래 보상의 기댓값 $J(\theta) = \mathbb{E}_{\tau \sim p(\tau; \theta)} [r(\tau)]$

$$= \int_{\tau} r(\tau) p(\tau; \theta) d\tau$$



미분

$$\nabla_{\theta} J(\theta) = \int_{\tau} r(\tau) \nabla_{\theta} p(\tau; \theta) d\tau \quad (\text{계산 불가능})$$



Monte carlo sampling

$$\nabla_{\theta} p(\tau; \theta) = p(\tau; \theta) \frac{\nabla_{\theta} p(\tau; \theta)}{p(\tau; \theta)} = p(\tau; \theta) \nabla_{\theta} \log p(\tau; \theta)$$

$$\begin{aligned} \nabla_{\theta} J(\theta) &= \int_{\tau} (r(\tau) \nabla_{\theta} \log p(\tau; \theta)) p(\tau; \theta) d\tau \\ &= \mathbb{E}_{\tau \sim p(\tau; \theta)} [r(\tau) \nabla_{\theta} \log p(\tau; \theta)] \end{aligned}$$

02 REINFORCE algorithm

$$p(\tau; \theta) = \prod_{t \geq 0} p(s_{t+1} | s_t, a_t) \pi_{\theta}(a_t | s_t)$$

↓ log

$$\log p(\tau; \theta) = \sum_{t \geq 0} \log p(s_{t+1} | s_t, a_t) + \log \pi_{\theta}(a_t | s_t)$$

↓ 미분

$$\nabla_{\theta} \log p(\tau; \theta) = \sum_{t \geq 0} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$$

전이 확률을 몰라도 $J(\theta)$ 미분 값 계산 가능!

$$\nabla_{\theta} J(\theta) \approx \sum_{t \geq 0} r(\tau) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$$

02 REINFORCE algorithm

Gradient estimator: $\nabla_{\theta} J(\theta) \approx \sum_{t \geq 0} r(\tau) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$

어떤 경로로부터 얻은 보상 $r(\tau)$ 이 크다면, 그 행동들을 할 확률이 높아짐
어떤 경로로부터 얻은 보상 $r(\tau)$ 이 작다면, 그 행동들을 할 확률이 낮아짐

어떤 경로가 좋다는 것은 그 경로에 포함되는 모든 행동이 좋았다는 것을 의미
-> 기댓값에 의해서 모두 averages out 됨
-> 구체적으로 어떤 행동이 좋은지 알 수 없음

02 REINFORCE algorithm

Gradient estimator: $\nabla_{\theta} J(\theta) \approx \sum_{t \geq 0} r(\tau) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$

어떤 경로로부터 얻은 보상 $r(\tau)$ 이 크다면, 그 행동들을 할 확률이 높아짐
어떤 경로로부터 얻은 보상 $r(\tau)$ 이 작다면, 그 행동들을 할 확률이 낮아짐

어떤 경로가 좋다는 것은 그 경로에 포함되는 모든 행동이 좋았다는 것을 의미
-> 기댓값에 의해서 모두 averages out 됨
-> 구체적으로 어떤 행동이 좋은지 알 수 없음

문제점: 높은 분산(high variance)
-> 분산을 낮추고 충분한 샘플링을 통해 estimator의 성능을 높여야 함

03 Variance reduction

분산을 줄이는 방법

1. 특정 상태에서부터 받을 미래 보상만을 고려하여 어떤 행동을 취할 확률을 키우는 방법

Gradient estimator:
$$\nabla_{\theta} J(\theta) \approx \sum_{t \geq 0} r(\tau) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$$

First idea: Push up probabilities of an action seen, only by the cumulative future reward from that state

$$\nabla_{\theta} J(\theta) \approx \sum_{t \geq 0} \left(\sum_{t' \geq t} r_{t'} \right) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$$

2. 지연된 보상에 대해 할인률을 적용하는 방법

Second idea: Use discount factor γ to ignore delayed effects

$$\nabla_{\theta} J(\theta) \approx \sum_{t \geq 0} \left(\sum_{t' \geq t} \gamma^{t'-t} r_{t'} \right) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$$

03 Variance reduction

3. Baseline

중요한 것은 실제로 얻은 보상이 얻을 것이라고 예상했던 것보다 좋은지 아닌지를 판단하는 것
→ Baseline function을 사용, 상태를 이용하는 방법!

Idea: Introduce a baseline function dependent on the state.

Concretely, estimator is now:

$$\nabla_{\theta} J(\theta) \approx \sum_{t \geq 0} \left(\sum_{t' \geq t} \gamma^{t'-t} r_{t'} - b(s_t) \right) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$$

Baseline function

- 해당 상태에서 얼마만큼의 보상을 원하는지 설명해주는 함수

03 Variance reduction

Baseline을 선택하는 방법

$$\nabla_{\theta} J(\theta) \approx \sum_{t \geq 0} \left(\sum_{t' \geq t} \gamma^{t'-t} r_{t'} - b(s_t) \right) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$$

1. 단순한 baseline

- 지금까지 경험했던 보상들에 대해서 moving average를 계산하는 방법
- 이런 variance reduction 방법이 “**vanilla REINFORCE**”
- 할인율을 적용, 미래에 받을 보상의 누적합을 계산, 단순한 baseline 추가

03 Variance reduction

$$\nabla_{\theta} J(\theta) \approx \sum_{t \geq 0} \left(\sum_{t' \geq t} \gamma^{t'-t} r_{t'} - b(s_t) \right) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$$

2. 더 좋은 baseline

- 우리는 어떤 행동이 그 상태에서의 기댓값보다 좋은 경우에 그 행동을 수행할 확률이 커지기를 원함
→ 여기서 Q-function과 value function을 이용

$Q^{\pi}(s_t, a_t) - V^{\pi}(s_t)$ 이 클수록 현재 행동이 좋음을 의미

$$\nabla_{\theta} J(\theta) \approx \sum_{t \geq 0} (Q^{\pi_{\theta}}(s_t, a_t) - V^{\pi_{\theta}}(s_t)) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$$

03 Variance reduction

$$\nabla_{\theta} J(\theta) \approx \sum_{t \geq 0} \left(\sum_{t' \geq t} \gamma^{t'-t} r_{t'} - b(s_t) \right) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$$

2. 더 좋은 baseline

- 우리는 어떤 행동이 그 상태에서의 기댓값보다 좋은 경우에 그 행동을 수행할 확률이 커지기를 원함
→ 여기서 Q-function과 value function을 이용

$Q^{\pi}(s_t, a_t) - V^{\pi}(s_t)$ 이 클수록 현재 행동이 좋음을 의미

이때 구체적인 Q-function과 Value function을 몰라도 됨
Q-learning과 **Policy gradient**를 이용!

04 Actor-Critic Algorithm

- Q-function과 Value function을 몰라도 **Q-learning**과 **Policy gradient**를 조합해서 training 시킬 수 있음
- Actor(Policy) – 우리가 어떤 행동을 할지 결정
- Critic(Q-function) – 그 행동이 얼마나 좋았는지, 또 어떻게 조절해야 하는지 알려줌
- Advantage function
 - 그 행동이 예상했던 것보다 얼마나 더 큰 보상을 주는지 알려주는 함수

$$A^{\pi}(s, a) = Q^{\pi}(s, a) - V^{\pi}(s)$$

04 Actor-Critic Algorithm

Initialize policy parameters θ , critic parameters ϕ

For iteration=1, 2 ... **do**

Sample m trajectories under the current policy

$\Delta\theta \leftarrow 0$

For $i=1, \dots, m$ **do**

For $t=1, \dots, T$ **do**

$$A_t = \sum_{t' \geq t} \gamma^{t'-t} r_{t'}^i - V_{\phi}(s_t^i)$$

$$\Delta\theta \leftarrow \Delta\theta + A_t \nabla_{\theta} \log(a_t^i | s_t^i)$$

$$\Delta\phi \leftarrow \sum_i \sum_t \nabla_{\phi} \|A_t^i\|^2$$

$$\theta \leftarrow \alpha \Delta\theta$$

$$\phi \leftarrow \beta \Delta\phi$$

End for

1. θ, ϕ 를 초기화 시킨다.
2. 현재의 policy를 기반으로 M 개의 경로를 샘플링한다.
3. Gradient를 계산한다. 각 경로마다 보상 함수를 계산하고 이용한다.
4. 보상함수를 이용해서 gradient estimator를 계산하고 이를 전부 누적시킨다.
5. ϕ 를 학습시키기 위해 가치 함수를 학습시킨다. 이는 보상 함수를 최소화시키는 것과 동일하므로, 가치 함수가 벨만 방정식(Bellman equation)에 근접하도록 학습시킨다.
6. 앞선 단계를 계속 반복한다.


Recurrent Attention Model(RAM)

- Hard attention 기법
- Image classification task에서 이미지의 glimpse만 가지고 예측
- 이미지 전체가 아닌 지역적인 부분만을 봄

Objective: Image Classification

Take a sequence of “glimpses” selectively focusing on regions of the image, to predict class

- Inspiration from human perception and eye movements
- Saves computational resources => scalability
- Able to ignore clutter / irrelevant parts of image



Recurrent Attention Model(RAM)


- Hard attention 기법
- Image classification task에서 이미지의 glimpse만 가지고 예측
- 이미지 전체가 아닌 지역적인 부분만을 봄

Objective: Image Classification

Take a sequence of “glimpses” selectively focusing on regions of the image, to predict class

- Inspiration from human perception and eye movements
- Saves computational resources => scalability
- Able to ignore clutter / irrelevant parts of image

이 문제를 강화학습으로 풀어보자!



Recurrent Attention Model(RAM)

Objective: Image Classification


Take a sequence of “glimpses” selectively focusing on regions of the image, to predict class

- Inspiration from human perception and eye movements
- Saves computational resources => scalability
- Able to ignore clutter / irrelevant parts of image

State: Glimpses seen so far

Action: (x,y) coordinates (center of glimpse) of where to look next in image

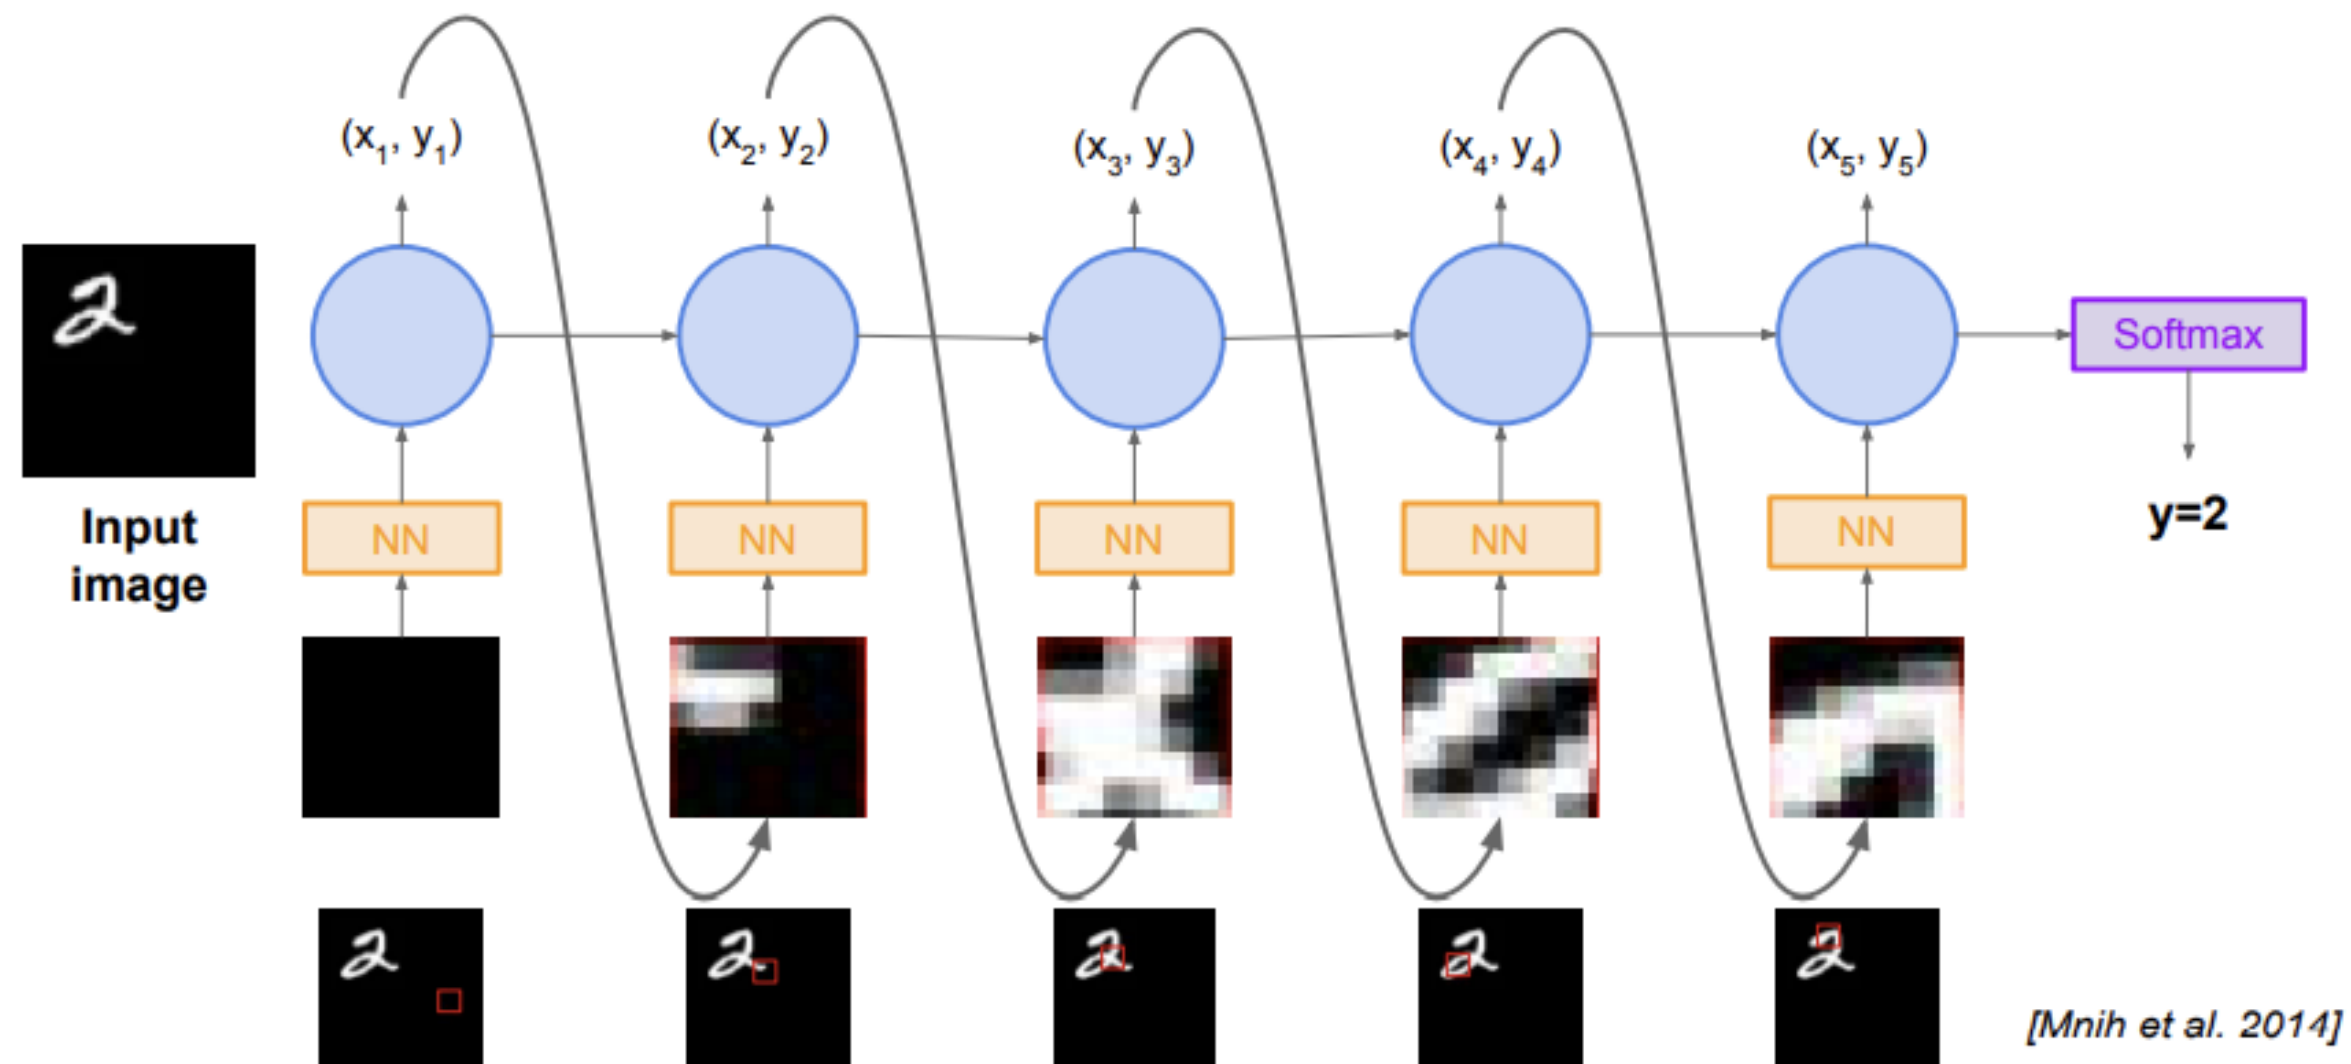
Reward: 1 at the final timestep if image correctly classified, 0 otherwise



- 상태는 지금까지 관찰한 glimpses
- 행동은 다음에 어떤 부분을 볼 것인지를 결정하는 것
- 보상은 classification의 성공 유무
 - > 어떻게 glimpses를 얻어낼 것인지를 정책을 통해 학습!

상태를 모델링하기 위해 RNN을 이용

REINFORCE in action: Recurrent Attention Model (RAM)



[Mnih et al. 2014]


More policy gradients: AlphaGo

Overview:

- Mix of supervised learning and reinforcement learning
- Mix of old methods (Monte Carlo Tree Search) and recent ones (deep RL)

How to beat the Go world champion:

- Featurize the board (stone color, move legality, bias, ...)
- Initialize policy network with supervised training from professional go games, then continue training using policy gradient (play against itself from random previous iterations, +1 / -1 reward for winning / losing)
- Also learn value network (critic)
- Finally, combine combine policy and value networks in a Monte Carlo Tree Search algorithm to select actions by lookahead search



[Silver et al.,
Nature 2016]

This image is CC0 public domain

THANK YOU

